

Unsupervised Data Quality Monitoring

THREE APPROACHES TO DETERMINE BATCH ACCEPTABILITY

Sheetal Laad, Micaela Flores, Ziv Schwartz, Kelly Sooch
Data Engineering for Machine Learning
NYU Center For Data Science

Introduction & Background

Task

Given a chronological series dataset batches, decide whether the next batch has acceptable data quality

Goal

Create and compare different methodologies that can discern which data batches are sufficiently clean by only training the model on a sample of clean batches

Motivation

When working with large amounts of data, it becomes too difficult to manually check if each input data batch maintains its integrity. Providing a dynamic data auditing system that performs unsupervised data quality monitoring given a chronological set of clean and dirty data batches to combat this issue.



Data

Data is comprised of Flights data and Facebook Posts

Facebook: 53 clean csv files and 53 dirty csv files

- 'line', 'page',
- 'week', 'num_likes',
- 'domain', 'outlet',
- 'title', 'description',
- 'contenttype', 'image',
- 'url', 'text',
- 'id', 'right_of_center'

Flights: 31 clean csv files and 31 dirty csv files

- 'RowId', 'Source',
- 'Flight', 'ScheduledDeparture',
- 'ActualDeparture', 'DepartureGate'
- 'ScheduledArrival', 'ActualArrival',
- 'ArrivalGate', 'for_key', 'date'



Data

Example Flight Data (same batch)

Clean



RowId	Source	Flight	ScheduledDeparture	ActualDeparture	DepartureGate	ScheduledArrival	ActualArrival	ArrivalGate	for_key	date
0	ua	UA-2708-EWR-CLT	Thu_ Dec 1 2:55 PM	Thu_ Dec 1 2:55 PM	A37	Thu_ Dec 1 4:53 PM	Thu_ Dec 1 4:44 PM	C5	Thu_ Dec 1 2:55 PM	734472
1	airtravelcenter	UA-2708-EWR-CLT	Thu_ Dec 1 2:55 PM	Thu_ Dec 1 2:55 PM	A37	Thu_ Dec 1 4:53 PM	Thu_ Dec 1 4:44 PM	C5	Thu_ Dec 1 2:55 PM	734472
2	myrateplan	UA-2708-EWR-CLT	Thu_ Dec 1 2:55 PM	Thu_ Dec 1 2:55 PM	A37	Thu_ Dec 1 4:53 PM	Thu_ Dec 1 4:44 PM	C5	Thu_ Dec 1 2:55 PM	734472
3	helloflight	UA-2708-EWR-CLT	Thu_ Dec 1 2:55 PM	Thu_ Dec 1 2:55 PM	A37	Thu_ Dec 1 4:53 PM	Thu_ Dec 1 4:44 PM	C5	Thu_ Dec 1 2:55 PM	734472
4	flytecomm	UA-2708-EWR-CLT	Thu_ Dec 1 2:55 PM	Thu_ Dec 1 2:55 PM	A37	Thu_ Dec 1 4:53 PM	Thu_ Dec 1 4:44 PM	C5	Thu_ Dec 1 2:55 PM	734472

Dirty



RowId	Source	Flight	ScheduledDeparture	ActualDeparture	DepartureGate	ScheduledArrival	ActualArrival	ArrivalGate	date
0	ua	UA-2708-EWR-CLT	Thu_ Dec 1 2:55 PM	Thu_ Dec 1 2:55 PM	A37	Thu_ Dec 1 4:53 PM	Thu_ Dec 1 4:44 PM	C5	734472
1	airtravelcenter	UA-2708-EWR-CLT	NaN	12/1/11 3:04 PM (-05:00)	NaN	NaN	12/1/11 4:22 PM (-05:00)	NaN	734472
2	myrateplan	UA-2708-EWR-CLT	NaN	12/1/11 3:04 PM (-05:00)	NaN	NaN	12/1/11 4:22 PM (-05:00)	NaN	734472
3	helloflight	UA-2708-EWR-CLT	NaN	12/1/11 3:04 PM (-05:00)	NaN	NaN	12/1/11 4:22 PM (-05:00)	NaN	734472
4	flytecomm	UA-2708-EWR-CLT	NaN	12/1/11 3:04 PM (-05:00)	NaN	NaN	12/1/11 4:22 PM (-05:00)	NaN	734472

Training Methodology

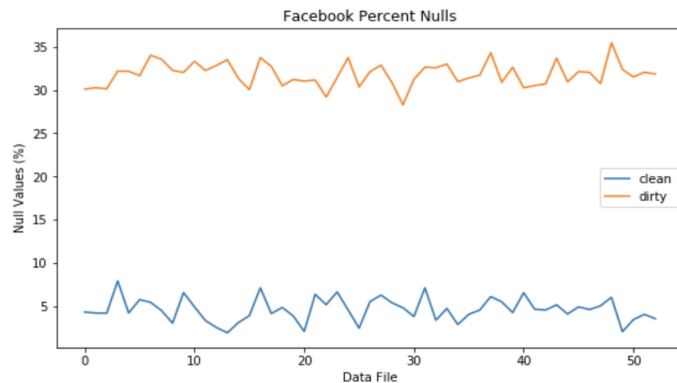
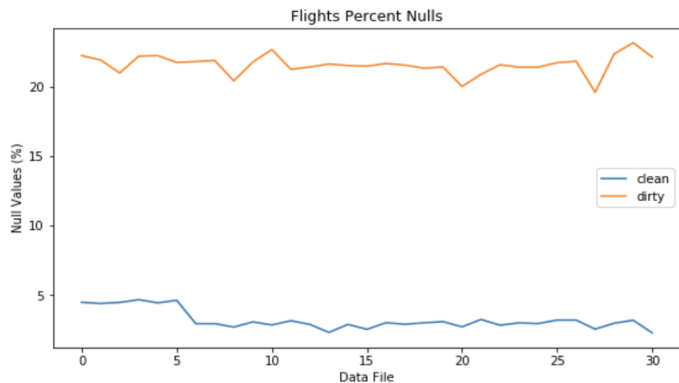
There were two different ways in which our group decided to handle the batch size used to train our models:

- **Increasing:**
 - Predetermined starting batch size is set and for every new batch that is read in, it is included as part of the whole training set.
 - Example: Starting batch size of 3, the first example would get trained on clean batches 1 to 3 and batch 4 will be tested. For the second example, clean batches 1 to 4 would be used to train and batch 5 will be tested.
- **Rolling:**
 - Predetermined batch size is used for the training criterion and moves from subsequent batch to batch
 - Example: Batch size set to 3, first example trained on batches 1 to 3 and batch 4 is tested. Second example is trained on batches 2 to 4, batch 5 tested.



Criterion Approach 1: Baseline Methodology

Percentage of null values per data batch



Criterion Explained:

- Percentage of null values in each batch of the training data
- If the percentage of nulls in the test batch is not within the minimum and maximum percentage of nulls of the training batches, the batch is classified as not acceptable

Criterion Approach 2: Range Methodology


- Uniqueness = count of distinct values by field
- Completeness = count of records that have a missing value for that field / count of records that don't have a missing value by field

Aggregated the mean values across all batches for each given column:

Uniqueness	Clean-----		Dirty-----	
	RowId	2381.290323	RowId	2381.290323
	Source	37.612903	Source	37.612903
	Flight	93.483871	Flight	93.483871
	ScheduledDeparture	82.935484	ScheduledDeparture	508.064516
	ActualDeparture	86.516129	ActualDeparture	662.903226
	DepartureGate	67.000000	DepartureGate	147.225806
	ScheduledArrival	90.354839	ScheduledArrival	588.225806
	ActualArrival	86.322581	ActualArrival	728.419355
	ArrivalGate	66.838710	ArrivalGate	146.225806
	for_key	82.935484	date	1.000000
	date	1.000000		
Completeness	Clean-----		Dirty-----	
	RowId	0.000000	RowId	0.000000
	Source	0.000000	Source	0.000000
	Flight	0.000000	Flight	0.000000
	ScheduledDeparture	0.000000	ScheduledDeparture	0.476794
	ActualDeparture	0.014968	ActualDeparture	0.190283
	DepartureGate	0.188292	DepartureGate	1.503791
	ScheduledArrival	0.000000	ScheduledArrival	0.464570
	ActualArrival	0.025478	ActualArrival	0.197457
	ArrivalGate	0.188728	ArrivalGate	1.488697
	date	0.000000	date	0.000000

Criterion Approach 2: Range Methodology

Completeness and Uniqueness by batch



Criterion Explained:

- If 80% of fields are within ± 2 stds of completeness mean and of uniqueness mean, batch is determined to be acceptable

Criterion Approach 3: TensorFlow Data Validation

Example of anomalies for a field:


Anomaly short description		Anomaly long description
Feature name		
'DepartureGate'	Unexpected string values	Examples contain values missing from the schema: - (~2%), 134 (~1%), 15? (<1%), 27/T2 (<1%), 29 (<1%), 3/8 (<1%), 33/8 (<1%), 37 (<1%), 40/8 (<1%), 41 (<1%), 41/8 (<1%), 44E (<1%), 44F/T4 (<1%), 45/T4 (<1%), 48B (<1%), 48B/T4 (<1%), 6/3 (<1%), 60 (<1%), 62 (<1%), 64 (<1%), 7 (<1%), 70A (<1%), 70B (~1%), 8/8 (<1%), 88 (<1%), 92 (<1%), A1 (<1%), A14 (<1%), A14\$ (<1%), A7/A (<1%), B22 (<1%), B3 (<1%), B8/B (<1%), C-33 (<1%), C11 (~3%), C134 (<1%), C16 (~1%), C16/C (<1%), C19 (<1%), C19/C (<1%), C30/C (<1%), C40 (<1%), C41 (<1%), C8/C (<1%), C88 (<1%), CHK (<1%), D20/D (<1%), D28 (~1%), D28/D (<1%), D30 (~1%), D32 (<1%), D4 (<1%), D48 (<1%), D75 (<1%), E1 (<1%), E2 (<1%), E8 (<1%), G14 (<1%), G19A/T3 (<1%), G1B/T3 (<1%), G3/T3 (<1%), G9/T3 (<1%), G92 (<1%), H15/T3 (<1%), K8/T3 (<1%), K9/T3 (<1%), Main Term - C25 (<1%), Not provided by airline (~1%), Term C - C134 (<1%), Term C - C33 (<1%), Term C - E1 (<1%), Terminal 2 (~1%), Terminal 2 - 57 (<1%), Terminal 3 (~1%), Terminal 3 - E3 (<1%), Terminal 4 (<1%), Terminal 4 - 44F (<1%), Terminal 4 - 45 (<1%), Terminal 4 - 48B (<1%), Terminal 6 (<1%), Terminal 8 - 3 (<1%), Terminal 8 - 33 (<1%), Terminal 8 - 40 (<1%), Terminal 8 - 41 (<1%), Terminal 8 - 8 (<1%), Terminal A - 15 (<1%), Terminal D - D1 (<1%), Terminal D - D36 (<1%), Terminal D - D40 (<1%), Terminal D - D48 (<1%).

Criterion Explained:





- TensorFlow Data Validation library to create an acceptable schema from the training batches and measure the number anomalies
- Anomalies are counted before and after schema adjustment, which consists of relaxing certain columns' domain mass requirements in the trained schema
 - Flights: 'DepartureGate' and 'ArrivalGate' were both reduced to 80% adherence
 - Facebook: 'outlet' and 'domain' were relaxed to 65% and 60% adherence, respectively
- Then number of anomalies is calculated again, using the same method from before. If the schema adjustment reduces the number of anomalies, then the test batch is considered acceptable

Evaluation/Results

Example of experiment results:



Final Results:



Discussion

The results of this project provides great insight into data quality monitoring in a production environment.

- There exists an optimal batch size
- Rolling training methods: faster runtimes, give less importance on older information (helpful in production when a field starts getting recorded different, when data seems to be changing, to capture seasonal trends)
- Range method: faster runtimes, more intuitive and generalizable

Some of the challenges that we faced:

- Inherent differences in production batches (Flights data)
- *TensorFlow Data Validation's* very long runtime
- *TensorFlow Data Validation* is not very generalizable

Future work for this topic could include:

- Looking more at the data values themselves
 - Involving more analysis of data trends (seasonal, yearly, etc.)
- 